

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
import io
```

```
df=pd.read_csv('Churn_Modelling.csv')
df
```



	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
...
9995	9996	15606229	Obijiaku	771	France	Male	39
9996	9997	15569892	Johnstone	516	France	Male	35
9997	9998	15584532	Liu	709	France	Female	36
9998	9999	15682355	Sabbatini	772	Germany	Male	42
9999	10000	15628319	Walker	792	France	Female	28

10000 rows × 14 columns

Next
steps:

[Generate code with df](#)



[View recommended plots](#)

[New interactive sheet](#)

```
df.shape
```



(10000, 14)

```
df=df.drop(['RowNumber','CustomerId','Surname'],axis=1)
df.head()
```



	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	Ha
0	619	France	Female	42	2	0.00	1	
1	608	Spain	Female	41	1	83807.86	1	
2	502	France	Female	42	8	159660.80	3	
3	699	France	Female	39	1	0.00	2	

4 850 Spain Female 43 2 125510.82 1

Next
steps:

[Generate code
with](#) `df`



[View recommended
plots](#)

[New interactive
sheet](#)

```
df.isna().any()  
df.isna().sum()
```



	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0

dtype: int64

```
print(df.shape)  
df.info()
```



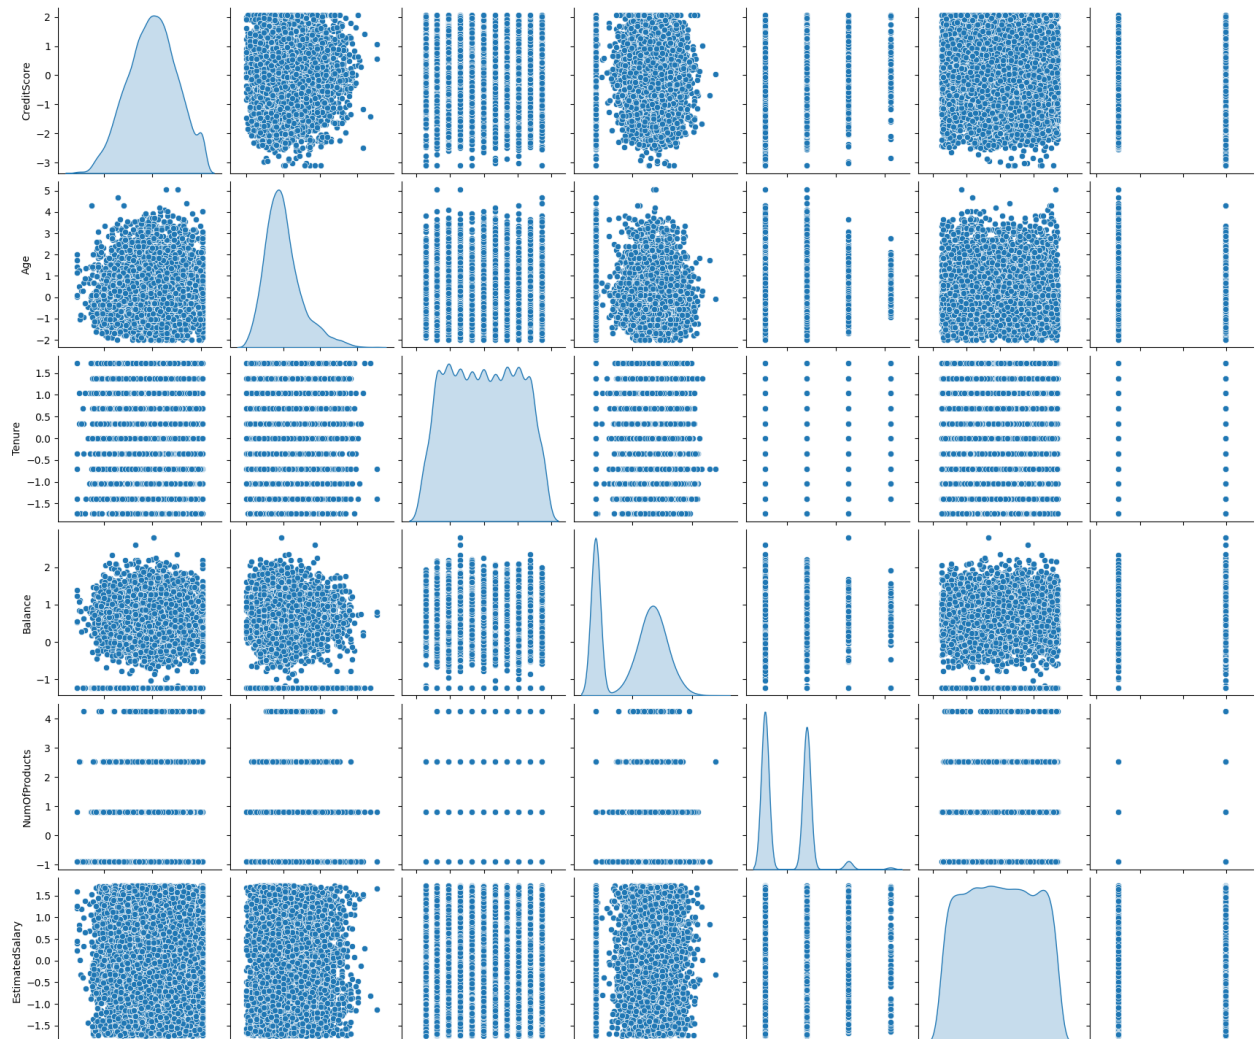
```
(10000, 11)  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10000 entries, 0 to 9999  
Data columns (total 11 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   CreditScore            10000 non-null  int64  
1   Geography              10000 non-null  object  
2   Gender                 10000 non-null  object  
3   Age                   10000 non-null  int64  
4   Tenure                 10000 non-null  int64  
5   Balance                10000 non-null  float64  
6   NumOfProducts          10000 non-null  int64  
7   HasCrCard              10000 non-null  int64  
8   IsActiveMember         10000 non-null  int64  
9   EstimatedSalary        10000 non-null  float64  
10  Exited                  10000 non-null  int64  
dtypes: float64(2), int64(7), object(2)  
memory usage: 859.5+ KB
```


```
df.describe()
```

	CreditScore	Age	Tenure	Balance	NumOfProducts
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	650.528800	38.921800	5.012800	76485.889288	1.530200
std	96.653299	10.487806	2.892174	62397.405202	0.581600
min	350.000000	18.000000	0.000000	0.000000	1.000000
25%	584.000000	32.000000	3.000000	0.000000	1.000000
50%	652.000000	37.000000	5.000000	97198.540000	1.000000
75%	718.000000	44.000000	7.000000	127644.240000	2.000000
max	850.000000	92.000000	10.000000	250898.090000	4.000000

```
scaler=StandardScaler()  
subset=df.drop(['Geography', 'Gender', 'HasCrCard', 'IsActiveMember'],axis=1)  
scaled=scaler.fit_transform(subset)  
scaled_df=pd.DataFrame(scaled,columns=subset.columns)  
sns.pairplot(scaled_df,diag_kind='kde')
```


<seaborn.axisgrid.PairGrid at 0x7b66931f94b0>








```
sns.heatmap(scaled_df.corr(),annot=True,cmap='rainbow')
```

<Axes: >



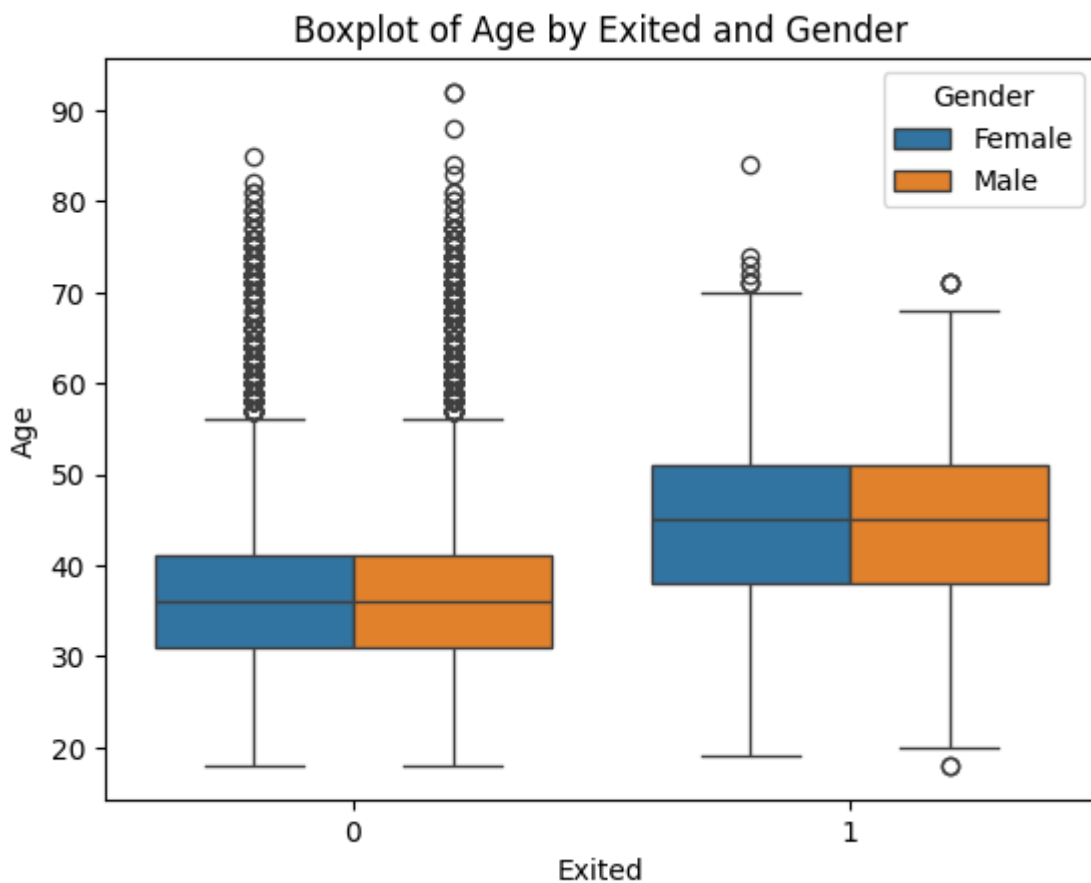
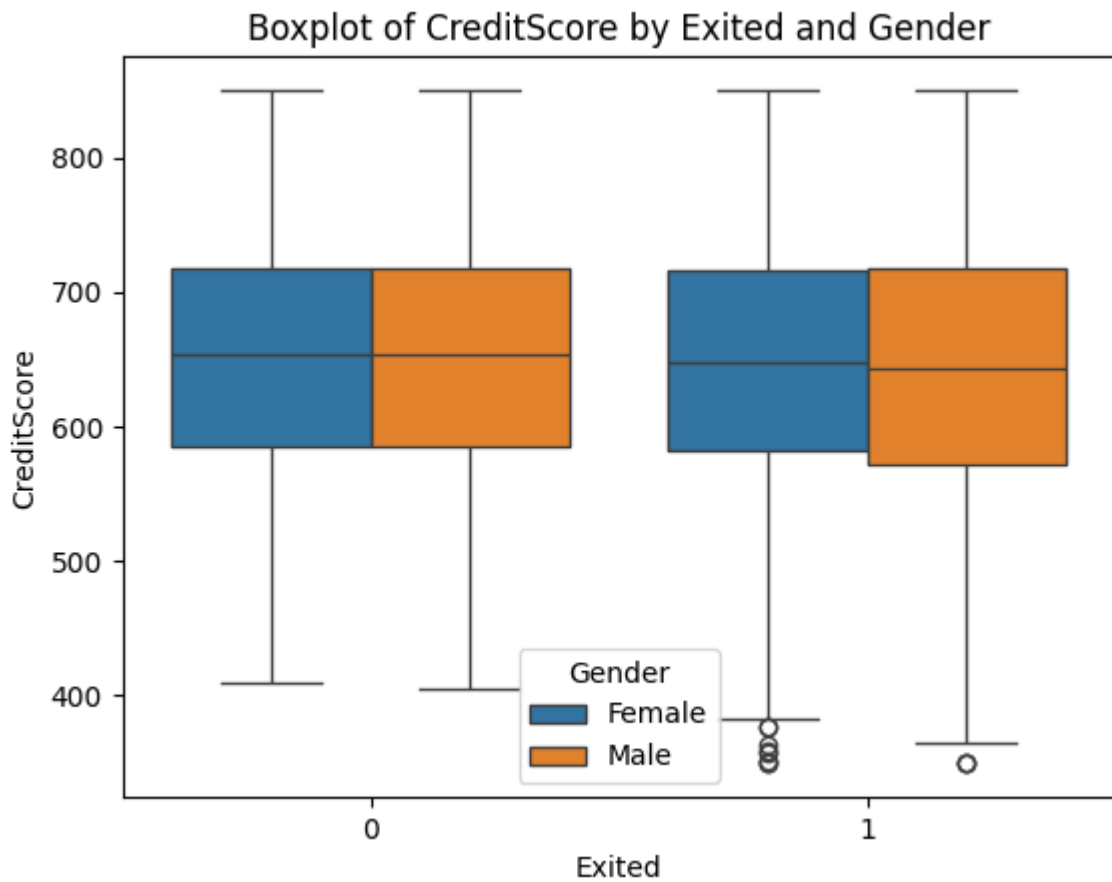
```
sns.countplot(x='Geography',data=df,hue='Exited')
plt.show()
sns.countplot(x='Gender',data=df,hue='Exited')
plt.show()
sns.countplot(x='HasCrCard',data=df,hue='Exited')
plt.show()
sns.countplot(x='IsActiveMember',data=df,hue='Exited')
plt.show()
```



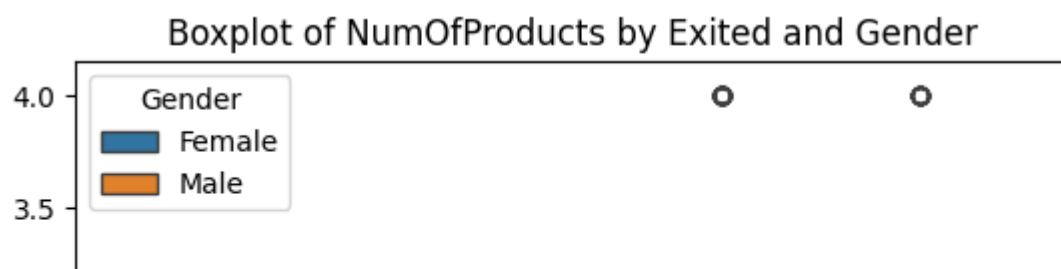
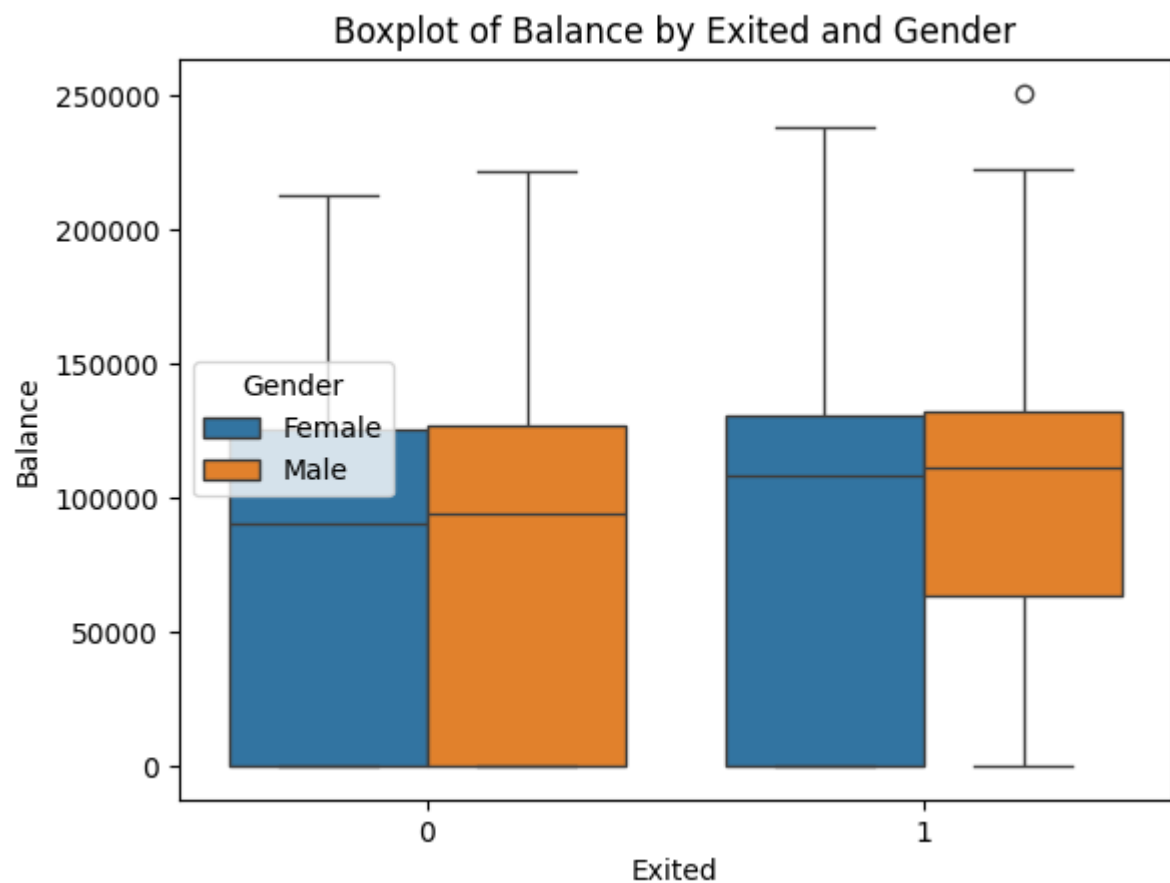



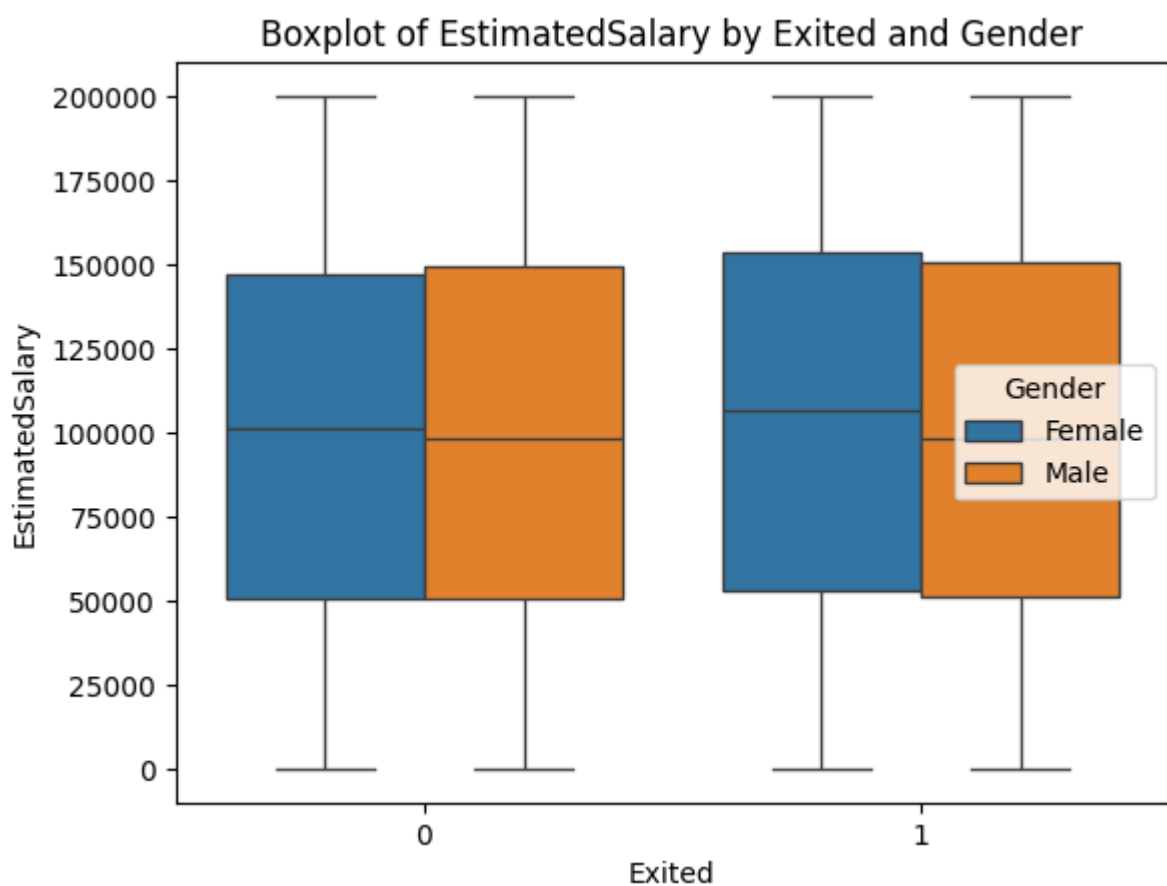

```
for i in subset.columns:  
    sns.boxplot(x=df['Exited'], y=df[i], hue=df['Gender'])  
    plt.title(f'Boxplot of {i} by Exited and Gender')
```

```
plt.title('Boxplot of {1} by Exited and Gender')  
plt.show()
```



Boxplot of Tenure by Exited and Gender





```
X=df.drop('Exited',axis=1)
y=df.pop('Exited')
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.10,random_state=42)
X_train,X_val,y_train,y_val=train_test_split(X_train,y_train,test_size=0.10,random_state=42)
print("X_train size is {}".format(X_train.shape[0]))
print("X_val size is {}".format(X_val.shape[0]))
print("X_test size is {}".format(X_test.shape[0]))
```

```
X_train size is 8100
```

```
X_val size is 900
X_test size is 1000
```

```
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
num_cols=['CreditScore','Age','Tenure','Balance','NumOfProducts','EstimatedSal
num_subset=scaler.fit_transform(X_train[num_cols])
X_train_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_train_num_df['Geography']=list(X_train['Geography'])
X_train_num_df['Gender']=list(X_train['Gender'])
X_train_num_df['HasCrCard']=list(X_train['HasCrCard'])
X_train_num_df['IsActiveMember']=list(X_train['IsActiveMember'])
X_train_num_df.head()
num_subset=scaler.fit_transform(X_val[num_cols])
X_val_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_val_num_df['Geography']=list(X_val['Geography'])
X_val_num_df['Gender']=list(X_val['Gender'])
X_val_num_df['HasCrCard']=list(X_val['HasCrCard'])
X_val_num_df['IsActiveMember']=list(X_val['IsActiveMember'])
num_subset=scaler.fit_transform(X_test[num_cols])
X_test_num_df=pd.DataFrame(num_subset,columns=num_cols)
X_test_num_df['Geography']=list(X_test['Geography'])
X_test_num_df['Gender']=list(X_test['Gender'])
X_test_num_df['HasCrCard']=list(X_test['HasCrCard'])
X_test_num_df['IsActiveMember']=list(X_test['IsActiveMember'])

X_train_num_df=pd.get_dummies(X_train_num_df,columns=['Geography','Gender'])
X_test_num_df=pd.get_dummies(X_test_num_df,columns=['Geography','Gender'])
X_val_num_df=pd.get_dummies(X_val_num_df,columns=['Geography','Gender'])
X_train_num_df.head()
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	EstimatedSala
0	-1.178587	-1.041960	-1.732257	0.198686	0.820905	1.5603
1	-0.380169	-1.326982	1.730718	-0.022020	-0.907991	-0.7135
2	-0.349062	1.808258	-0.693364	0.681178	0.820905	-1.1265
3	0.625629	2.378302	-0.347067	-1.229191	0.820905	-1.6827
4	-0.203895	-1.136967	1.730718	0.924256	-0.907991	1.3325

Next
steps:

[Generate code with X_train_num_df](#)

 recommended

[New interactive sheet](#)

```
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense

model=Sequential()
```

```

model.add(Dense(7,activation='relu'))
model.add(Dense(10,activation='relu'))
model.add(Dense(1,activation='sigmoid'))

import tensorflow as tf
optimizer=tf.keras.optimizers.Adam(0.01)
model.compile(loss='binary_crossentropy',optimizer=optimizer,metrics=['accuracy'])

model.fit(X_train_num_df,y_train,epochs=100,batch_size=10,verbose=1)

```

Epoch 1/100
810/810 ————— 3s 1ms/step - accuracy: 0.8094 - loss: 0.4333
Epoch 2/100
810/810 ————— 1s 1ms/step - accuracy: 0.8595 - loss: 0.3542
Epoch 3/100
810/810 ————— 1s 1ms/step - accuracy: 0.8634 - loss: 0.3406
Epoch 4/100
810/810 ————— 2s 2ms/step - accuracy: 0.8575 - loss: 0.3501
Epoch 5/100
810/810 ————— 3s 2ms/step - accuracy: 0.8562 - loss: 0.3511
Epoch 6/100
810/810 ————— 2s 1ms/step - accuracy: 0.8608 - loss: 0.3407
Epoch 7/100
810/810 ————— 1s 1ms/step - accuracy: 0.8603 - loss: 0.3399
Epoch 8/100
810/810 ————— 1s 1ms/step - accuracy: 0.8606 - loss: 0.3385
Epoch 9/100
810/810 ————— 1s 1ms/step - accuracy: 0.8590 - loss: 0.3385
Epoch 10/100
810/810 ————— 1s 1ms/step - accuracy: 0.8661 - loss: 0.3265
Epoch 11/100
810/810 ————— 1s 1ms/step - accuracy: 0.8639 - loss: 0.3331
Epoch 12/100
810/810 ————— 1s 1ms/step - accuracy: 0.8619 - loss: 0.3372
Epoch 13/100
810/810 ————— 2s 2ms/step - accuracy: 0.8683 - loss: 0.3298
Epoch 14/100
810/810 ————— 2s 2ms/step - accuracy: 0.8675 - loss: 0.3364
Epoch 15/100
810/810 ————— 1s 1ms/step - accuracy: 0.8654 - loss: 0.3223
Epoch 16/100
810/810 ————— 1s 1ms/step - accuracy: 0.8682 - loss: 0.3188
Epoch 17/100
810/810 ————— 1s 1ms/step - accuracy: 0.8646 - loss: 0.3332
Epoch 18/100
810/810 ————— 1s 1ms/step - accuracy: 0.8651 - loss: 0.3289
Epoch 19/100
810/810 ————— 1s 1ms/step - accuracy: 0.8636 - loss: 0.3340
Epoch 20/100
810/810 ————— 1s 1ms/step - accuracy: 0.8698 - loss: 0.3255
Epoch 21/100
810/810 ————— 1s 1ms/step - accuracy: 0.8751 - loss: 0.3104
Epoch 22/100
810/810 ————— 2s 2ms/step - accuracy: 0.8629 - loss: 0.3377
Epoch 23/100
810/810 ————— 2s 2ms/step - accuracy: 0.8677 - loss: 0.3260

```

810/810 ————— 2s 2ms/step - accuracy: 0.8677 - loss: 0.3209
Epoch 24/100
810/810 ————— 2s 1ms/step - accuracy: 0.8684 - loss: 0.3285
Epoch 25/100
810/810 ————— 1s 1ms/step - accuracy: 0.8667 - loss: 0.3321
Epoch 26/100
810/810 ————— 1s 1ms/step - accuracy: 0.8658 - loss: 0.3301
Epoch 27/100
810/810 ————— 1s 1ms/step - accuracy: 0.8589 - loss: 0.3379
Epoch 28/100
810/810 ————— 1s 1ms/step - accuracy: 0.8667 - loss: 0.3233
Epoch 29/100
810/810 ————— 1s 1ms/step - accuracy: 0.8646 - loss: 0.3410

```

```

y_pred_val=model.predict(X_val_num_df)
y_pred_val[y_pred_val>0.5]=1
y_pred_val[y_pred_val <0.5]=0

```

```

29/29 ————— 0s 2ms/step

```

```

y_pred_val=y_pred_val.tolist()
X_compare_val=X_val.copy()
X_compare_val['y_actual']=y_val
X_compare_val['y_pred']=y_pred_val
X_compare_val.head(10)

```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
340	642	Germany	Female	40	6	129502.49	2
8622	706	Germany	Male	36	9	58571.18	2
8401	535	Spain	Male	58	1	0.00	2
4338	714	Spain	Male	25	2	0.00	1
8915	606	France	Male	36	1	155655.46	1
2624	605	Spain	Female	29	3	116805.82	1
2234	720	France	Female	38	10	0.00	2
349	582	France	Male	39	5	0.00	2
3719	850	France	Female	62	1	124678.35	1
2171	526	Germany	Male	58	9	190298.89	2

Next
steps:

[Generate
code with](#) `X_compare_val`

[View recommended
plots](#)

[New interactive
sheet](#)

```

from sklearn.metrics import confusion_matrix
cm_val=confusion_matrix(y_val,y_pred_val)
cm_val

```

```
array([[672, 44],
       [ 89, 95]])
```

```
Accuracy=782/900
```

```
print("Accuracy of the Model on the Validation Data set is 86.89%")
```

```
Accuracy of the Model on the Validation Data set is 86.89%
```

```
loss1,accuracy1=model.evaluate(X_train_num_df,y_train,verbose=False)
```

```
loss2,accuracy2=model.evaluate(X_val_num_df,y_val,verbose=False)
```

```
print("Train Loss {}".format(loss1))
```

```
print("Train Accuracy {}".format(accuracy1))
```

```
print("Val Loss {}".format(loss2))
```

```
print("Val Accuracy {}".format(accuracy2))
```

```
Train Loss 0.3174085319042206
```

```
Train Accuracy 0.87086421251297
```

```
Val Loss 0.35060858726501465
```

```
Val Accuracy 0.852222204208374
```

```
from sklearn import metrics
```

```
y_pred_test=model.predict(X_test_num_df)
```

```
y_pred_test[y_pred_test>0.5]=1
```

```
y_pred_test[y_pred_test <0.5]=0
```

```
cm_test=metrics.confusion_matrix(y_test,y_pred_test)
```

```
cm_test
```

```
print("Test Confusion Matrix")
```

```
32/32 ————— 0s 2ms/step
```

```
Test Confusion Matrix
```

```
cm_test
```

```
array([[742, 52],
       [114, 92]])
```

```
loss3,accuracy3=model.evaluate(X_test_num_df,y_test,verbose=False)
```

```
print("Test Accuracy is {}".format(accuracy3))
```

```
print("Test loss is {}".format(loss3))
```

```
Test Accuracy is 0.8339999914169312
```

```
Test loss is 0.3814462721347809
```

```
Start coding or generate with AI.
```

